

Matrix Multiplication Compiler for a Custom 24-bit ISA

Aditya 22BAI1235¹ and Affan Ahmed 22BRS1141¹

Vellore Institute of Technology (VIT), Chennai, India
{aditya.2022a,affan.ahmed2022}@vitstudent.ac.in

Abstract. This document presents a project report for a custom compiler designed to translate matrix multiplication (both matrix–matrix and matrix–vector) operations written in C/C++ into a 24-bit Instruction Set Architecture (ISA). The report is structured as follows:

- **Analysis and Design of the Algorithm**
- **Source Code of the Solution with Explanations**
- **Output (Screenshots, Logs, and Generated Files)**

Placeholders are provided for images and output files where necessary.

Keywords: Compiler · Matrix Multiplication · Custom ISA · Code Generation

1 Analysis and Design of Algorithm

This section explains how the custom compiler transforms C/C++ code containing matrix multiplication into a 24-bit ISA.

1.1 Compiler Overview

The compiler follows a multi-stage pipeline (see Fig. 1):

1. **Parsing:** Extracts matrix dimensions and operation type (matrix–matrix or matrix–vector) from special markers in the source code (e.g., `// OPERATION: MM_MULT`).
2. **Semantic Analysis:** Validates the dimensions to ensure they match for multiplication.
3. **IR Optimization:** Performs minimal optimizations, such as counting multiply-accumulate (MAC) steps.
4. **Code Generation:** Produces machine instructions conforming to the 24-bit ISA.
5. **Memory Layout Computation:** Assigns base addresses for matrices (A, B, C) or vector V.

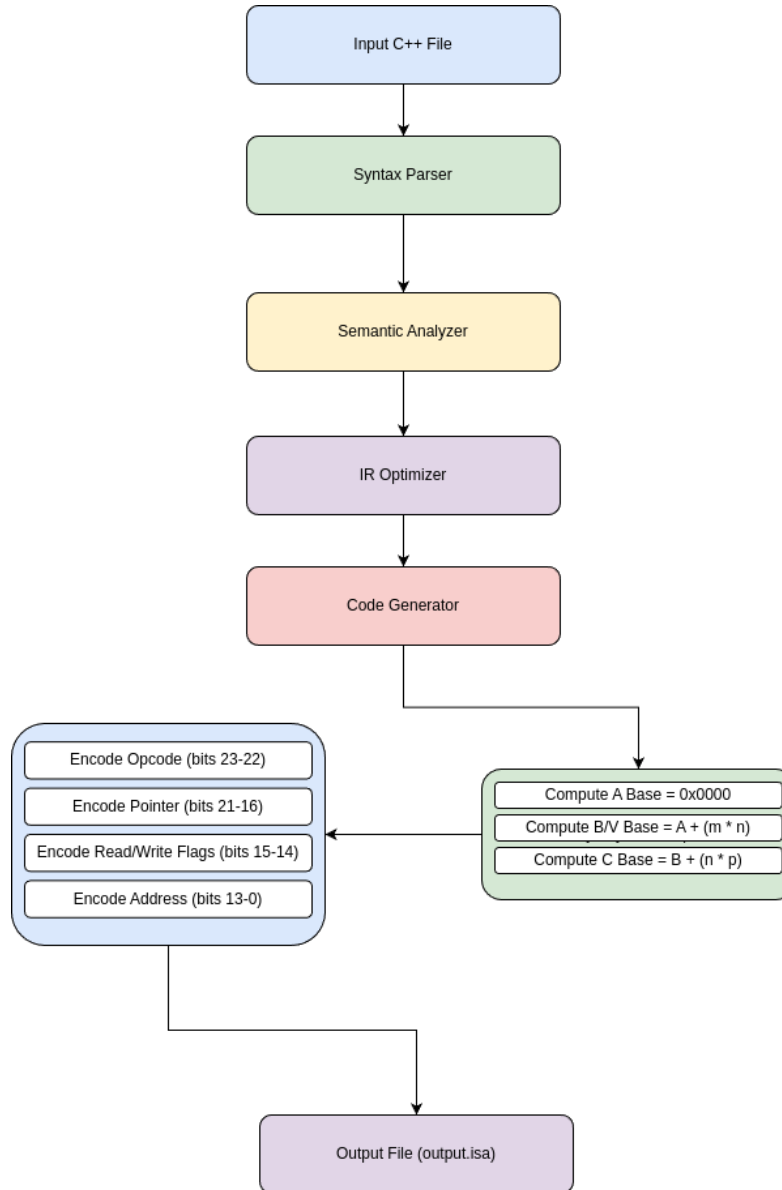


Fig. 1. High-level pipeline of the compiler operations.

1.2 24-bit ISA Format

Each instruction is encoded into 24 bits as follows:

- **Bits 23–22:** Opcode type (e.g., READ, PROG, EXE, END).
- **Bits 21–16:** Pointer field (used for microcode control or indexing).
- **Bit 15:** Read flag.
- **Bit 14:** Write flag.
- **Bits 13–0:** Address (lower 14 bits).

1.3 Memory Layout

Matrices are placed sequentially in memory:

- Matrix A starts at address 0x0000.
- Matrix B (or vector V) follows A.
- Matrix C is placed after B.

For example, if matrix A is of size $m \times n$, then B's base address is:

$$\text{base_B} = \text{base_A} + (m \times n)$$

and matrix C's base address is:

$$\text{base_C} = \text{base_B} + (n \times p)$$

where p is the number of columns in matrix B.

2 Source Code of the Solution

Below, the complete source code files are displayed in full. Each section includes an explanation of its purpose, functionality, and a small example excerpt to illustrate how it contributes to the custom ISA architecture.

2.1 compiler_main.c

Purpose: This file is the entry point of the compiler. It handles reading the input C++ file, invoking the parser to extract markers and dimensions, then sequentially calls the semantic analyzer, IR optimizer, and code generator to produce the final `output.isa` file.

Code:

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <string.h>
4 #include "syntax_parser.h"
5 #include "semantic_analyzer.h"
6 #include "ir_optimizer.h"

```

```

7 #include "isa_generator.h"
8
9 int main(int argc, char *argv[]) {
10     if (argc < 2) {
11         fprintf(stderr, "Usage: %s <source_file.cpp>\n", argv
12             [0]);
13         return 1;
14     }
15
16     // Open the input C++ source file.
17     FILE *fp = fopen(argv[1], "r");
18     if (!fp) {
19         perror("Error opening source file");
20         return 1;
21     }
22
23     // Determine the file size.
24     fseek(fp, 0, SEEK_END);
25     long file_size = ftell(fp);
26     rewind(fp);
27
28     // Allocate buffer and read the file contents.
29     char *source_code = malloc(file_size + 1);
30     if (!source_code) {
31         fprintf(stderr, "Memory allocation error\n");
32         fclose(fp);
33         return 1;
34     }
35     fread(source_code, 1, file_size, fp);
36     source_code[file_size] = '\0';
37     fclose(fp);
38
39     printf("Loaded Source Code:\n%s\n", source_code);
40
41     // Parse the source code to extract operation type and
42     // dimensions.
43     ParsedInfo parsed = parse_source(source_code);
44     printf("Parsed Info: Operation Type: %s, m=%d, n=%d, p=%d
45         \n",
46         (parsed.opType == MM_MULT) ? "Matrix-Matrix" : "
47         Matrix-Vector",
48         parsed.m, parsed.n, parsed.p);
49
50     // Perform semantic analysis to build the intermediate
51     // representation.
52     MatrixIR ir = perform_semantic_analysis(parsed);
53     printf("IR: Matrix A: %dx%d, Matrix B: %dx%d\n",
54         ir.rows_A, ir.cols_A, ir.rows_B, ir.cols_B);
55
56     // Optimize the IR.

```

```

52     OptimizedInfo opt = optimize_ir(&ir);
53     printf("Optimized Info: MAC steps=%d, partial mults=%d\n"
54           ,
55           opt.mac_steps, opt.partial_mults);
56     // Generate the ISA instruction sequence (this writes
57     //   output.isa).
58     generate_instruction_sequence(&ir, &opt);
59     free(source_code);
60     return 0;
61 }

```

Listing 1.1. compiler_main.c

Explanation: Lines 1–13 initialize the program and check for the required input file. Lines 15–36 read the entire C++ source file into a buffer. Lines 38–44 call the parser to extract the operation type and dimensions. Subsequent lines perform semantic analysis and optimization, and finally, the ISA generator is invoked to write the output file.

2.2 syntax_parser.c

Purpose: Extracts special markers from the C++ source code, such as `// OPERATION:` and `// DIMENSIONS:`, and creates a `ParsedInfo` structure containing the operation type and dimensions.

Code:

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4  #include "syntax_parser.h"
5
6  ParsedInfo parse_source(const char* source_code) {
7      ParsedInfo info;
8      // Set default values.
9      info.opType = MM_MULT; // default to matrix-matrix
10     info.m = 0;
11     info.n = 0;
12     info.p = 0;
13
14     // Look for the "OPERATION:" marker.
15     const char *op_marker = strstr(source_code, "OPERATION:");
16     ;
17     if (op_marker) {
18         op_marker += strlen("OPERATION:");
19         while (*op_marker == ' ' || *op_marker == '\t')
20             op_marker++;
21         if (strncmp(op_marker, "MV_MULT", 7) == 0)

```

```

20         info.opType = MV_MULT;
21     else
22         info.opType = MM_MULT;
23 }
24
25 // Look for the "DIMENSIONS:" marker.
26 const char *dim_marker = strstr(source_code, "DIMENSIONS:");
27 if (dim_marker) {
28     dim_marker += strlen("DIMENSIONS:");
29     char dims[128];
30     sscanf(dim_marker, "%127[^\n]", dims); // read until
31                                         newline
32
33     // Tokenize the dimensions string.
34     char *token = strtok(dims, " ");
35     while (token != NULL) {
36         if (strncmp(token, "m=", 2) == 0)
37             info.m = atoi(token + 2);
38         else if (strncmp(token, "n=", 2) == 0)
39             info.n = atoi(token + 2);
40         else if (strncmp(token, "p=", 2) == 0)
41             info.p = atoi(token + 2);
42         token = strtok(NULL, " ");
43     }
44
45     // For matrix-vector multiplication, set p to 1 if not
46     // provided.
47     if (info.opType == MV_MULT)
48         info.p = 1;
49
50     // Basic validation.
51     if (info.m <= 0 || info.n <= 0 || info.p <= 0) {
52         fprintf(stderr, "Invalid or missing dimensions in
53             source code.\n");
54         exit(EXIT_FAILURE);
55     }
56
57     return info;
58 }

```

Listing 1.2. syntax_parser.c

Explanation: The parser uses functions like `strstr` and `sscanf` to find and interpret the markers. For example, it defaults to a matrix-matrix operation unless the "MV_MULT" marker is found. It tokenizes the dimension string to extract values for m , n , and p , and ensures basic validation.

2.3 semantic_analyzer.c

Purpose: Validates the dimensions extracted by the parser (e.g., ensuring that the number of columns in matrix A equals the number of rows in matrix B) and builds the intermediate representation (IR) for the multiplication operation.

Code:

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include "semantic_analyzer.h"
4
5  MatrixIR perform_semantic_analysis(ParsedInfo info) {
6      MatrixIR ir;
7
8      if (info.opType == MM_MULT) {
9          // For matrix-matrix multiplication, assume:
10         // A is m x n, and B is n x p.
11         ir.rows_A = info.m;
12         ir.cols_A = info.n;
13         ir.rows_B = info.n;
14         ir.cols_B = info.p;
15     } else {
16         // For matrix-vector multiplication, assume:
17         // A is m x n and vector V is n x 1.
18         ir.rows_A = info.m;
19         ir.cols_A = info.n;
20         ir.rows_B = info.n;
21         ir.cols_B = 1;
22     }
23
24     // Validate that the inner dimensions match.
25     if (ir.cols_A != ir.rows_B) {
26         fprintf(stderr, "Dimension mismatch: A's columns (%d)
27             must equal B's rows (%d).\n", ir.cols_A, ir.
28             rows_B);
29         exit(EXIT_FAILURE);
30     }
31
32     return ir;
33 }
```

Listing 1.3. semantic_analyzer.c

Explanation: This module constructs a `MatrixIR` structure. For matrix-matrix multiplication, it assigns A as an $m \times n$ matrix and B as an $n \times p$ matrix. It also performs a dimension check and exits if a mismatch is detected.

2.4 ir_optimizer.c

Purpose: Calculates performance metrics such as the number of multiply-accumulate (MAC) operations required by the multiplication. It provides a basic optimization layer, even if minimal.

Code:

```

1 #include <stdio.h>
2 #include "ir_optimizer.h"
3
4 OptimizedInfo optimize_ir(MatrixIR *ir) {
5     OptimizedInfo opt;
6     // Calculate the number of Multiply-Accumulate (MAC)
7     // steps.
8     // For multiplication: MAC steps = rows_A * cols_B *
9     // cols_A
10    opt.mac_steps = ir->rows_A * ir->cols_B * ir->cols_A;
11    // For simplicity, set partial multiplications equal to
12    // MAC steps.
13    opt.partial_mults = opt.mac_steps;
14    return opt;
15 }
```

Listing 1.4. ir_optimizer.c

Explanation: The optimizer computes the total MAC steps as $m \times p \times n$ and sets the number of partial multiplications equal to this value. This information is passed on to the code generator for potential optimizations.

2.5 isa_encoding.c

Purpose: Implements the function to encode a 24-bit instruction word. The instruction is built by packing the opcode, pointer, read/write flags, and memory address into a single 24-bit integer.

Code:

```

1 #include "isa_encoding.h"
2 #include <stdio.h>
3 // Function to encode a 24-bit instruction.
4 unsigned int encode_instruction(int opcode_type, int pointer,
5     int read_flag, int write_flag, int address) {
6     unsigned int instruction = 0;
7     instruction |= ((opcode_type & 0x3) << 22);    // Bits
8     // 23-22.
9     instruction |= ((pointer & 0x3F) << 16);        // Bits
10    // 21-16.
11    instruction |= ((read_flag & 0x1) << 15);        // Bit
12    // 15.
13 }
```



```

9      instruction |= ((write_flag & 0x1) << 14);           // Bit
      14.
10     instruction |= (address & 0x3FFF);                   //
      Bits 13-0.
11     return instruction;
12 }
13
14 // Define a static lookup table for instruction opcodes.
15 static InstructionInfo instruction_table[] = {
16     {"READ", 0x00},    // Opcode for READ instruction.
17     {"PROG", 0x01},    // Opcode for PROG (start) instruction.
18     {"EXE", 0x02},     // Opcode for EXE (execute) instruction
      .
19     {"END", 0x03}      // Opcode for END instruction.
20 };
21
22 // Returns the number of instructions in the lookup table.
23 int get_instruction_table_size(void) {
24     return sizeof(instruction_table) / sizeof(
      instruction_table[0]);
25 }
26
27 // Returns a pointer to the InstructionInfo for the given
      index.
28 const InstructionInfo * get_instruction_info_by_index(int
      index) {
29     if (index < 0 || index >= get_instruction_table_size())
30         return NULL;
31     return &instruction_table[index];
32 }

```

Listing 1.5. isa_encoding.c

Explanation: The function `encode_instruction` takes five parameters and shifts/masks them appropriately. A static lookup table maps instruction names to opcodes for debugging and verification purposes.

2.6 isa_generator.c

Purpose: Generates the final instruction sequence. It uses the IR and memory layout to create READ instructions for matrix elements and control instructions (PROG, EXE, END) for each computed element of the result matrix. It then writes these instructions to `output.isa`.

Code:

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include "isa_generator.h"
4 #include "isa_encoding.h"

```

```

5 #include "target_memory.h"
6
7 // Generates a READ instruction using the computed memory
  layout.
8 static void generate_read_instruction(char operand, int
  index1, int index2, MatrixIR *ir, MemoryLayout *layout,
  FILE *fp) {
9     unsigned int address = 0;
10    if (operand == 'A') {
11        // For matrix A: address = base_A + (i * cols_A + k)
12        address = layout->base_A + (index1 * ir->cols_A +
  index2);
13    } else if (operand == 'B') {
14        // For matrix B (matrix-matrix): address = base_B + (
  k * cols_B + j)
15        address = layout->base_B + (index1 * ir->cols_B +
  index2);
16    } else if (operand == 'V') {
17        // For vector (matrix-vector): address = base_V +
  index1
18        address = layout->base_V + index1;
19    }
20    int opcode_type = 0; // READ opcode.
21    int pointer = 0;
22    int read_flag = 1;
23    int write_flag = 0;
24    unsigned int instruction = encode_instruction(opcode_type
  , pointer, read_flag, write_flag, address);
25    if (operand == 'A' || operand == 'B')
26        fprintf(fp, "  READ %c[%d][%d]: 0x%06X\n", operand,
  index1, index2, instruction);
27    else // For vector, index2 is not used.
28        fprintf(fp, "  READ %c[%d]: 0x%06X\n", operand,
  index1, instruction);
29 }
30
31 // Generates control instructions (PROG, EXE, END) for a
  result element.
32 static void generate_control_instructions(int i, int j,
  MatrixIR *ir, MemoryLayout *layout, FILE *fp) {
33     unsigned int address = 0;
34     if (ir->cols_B > 1) {
35         // Matrix-matrix multiplication: C[i][j] address.
36         address = layout->base_C + (i * ir->cols_B + j);
37     } else {
38         // Matrix-vector multiplication: C[i] address.
39         address = layout->base_C + i;
40     }
41     int pointer = 0;
42     int read_flag = 0;

```

```

43     int write_flag = 0;
44
45     unsigned int prog = encode_instruction(1, pointer,
46         read_flag, write_flag, address);
47     unsigned int exe = encode_instruction(2, pointer,
48         read_flag, write_flag, address);
49     unsigned int end = encode_instruction(3, pointer,
50         read_flag, write_flag, address);
51
52     fprintf(fp, "  PROG: 0x%06X\n", prog);
53     fprintf(fp, "  EXE:  0x%06X\n", exe);
54     fprintf(fp, "  END:  0x%06X\n", end);
55 }
56
57 void generate_instruction_sequence(MatrixIR *ir,
58     OptimizedInfo *opt) {
59     FILE *fp = fopen("output.isa", "w");
60     if (!fp) {
61         perror("Failed to open output.isa");
62         exit(EXIT_FAILURE);
63     }
64
65     // Write a pretty header.
66     fprintf(fp, "
67         =====\n");
68     fprintf(fp, "      Generated Instruction Sequence (24-bit\n");
69     fprintf(fp, "      ISA Format)      \n");
70     fprintf(fp, "
71         =====\n");
72     fprintf(fp, "      \n");
73
74     // Initialize a counter for the total number of
75     // instructions generated.
76     int total_instructions = 0;
77
78     // Compute memory layout based on IR.
79     MemoryLayout layout;
80     compute_memory_layout(ir, &layout);
81
82     // Generate instructions based on multiplication type.
83     if (ir->cols_B > 1) {
84         // Matrix-Matrix Multiplication.
85         int m = ir->rows_A;
86         int n = ir->cols_A; // Also equals ir->rows_B.
87         int p = ir->cols_B;
88         for (int i = 0; i < m; i++) {
89             for (int j = 0; j < p; j++) {
90                 fprintf(fp, "Processing C[%d][%d]\n", i, j);
91                 for (int k = 0; k < n; k++) {

```

```

83         // Read A[i][k]: computed from base_A.
84         generate_read_instruction('A', i, k, ir,
85             &layout, fp);
86         total_instructions++;
87         // Read B[k][j]: computed from base_B.
88         generate_read_instruction('B', k, j, ir,
89             &layout, fp);
90         total_instructions++;
91     }
92     // Control instructions for C[i][j]: PROG,
93     EXE, and END.
94     generate_control_instructions(i, j, ir, &
95         layout, fp);
96     total_instructions += 3; // Three control
97     instructions.
98     fprintf(fp, "\n");
99 }
100 }
101 } else {
102     // Matrix-Vector Multiplication.
103     int m = ir->rows_A;
104     int n = ir->cols_A;
105     for (int i = 0; i < m; i++) {
106         fprintf(fp, "Processing C[%d]\n", i);
107         for (int k = 0; k < n; k++) {
108             generate_read_instruction('A', i, k, ir, &
109                 layout, fp);
110             total_instructions++;
111             // For matrix-vector, use operand 'V' for
112             vector.
113             generate_read_instruction('V', k, 0, ir, &
114                 layout, fp);
115             total_instructions++;
116         }
117         generate_control_instructions(i, 0, ir, &layout,
118             fp);
119         total_instructions += 3; // Three control
120         instructions.
121         fprintf(fp, "\n");
122     }
123 }
124
125 int total_bits = total_instructions * 24;
126
127 fprintf(fp, "
128     =====\n");
129 fprintf(fp, "Total instructions generated: %d\n",
130     total_instructions);
131 fprintf(fp, "Total bits: %d\n", total_bits);

```

```

120     fprintf(fp, "
        =====\n
        n");
121
122     fclose(fp);
123     printf("Instruction sequence generated in output.isa\n");
124 }

```

Listing 1.6. isa_generator.c

Explanation: After computing the memory layout using `compute_memory_layout`, the generator loops over each result element. For each element, it generates READ instructions for corresponding elements of matrices A and B (or vector V), and then emits control instructions. A counter tallies the total instructions, and a summary is appended to the output file.

2.7 target_memory.c

Purpose: Computes base addresses for matrices A, B (or vector V), and C based on the provided dimensions. This ensures that every instruction referencing a matrix element uses the correct address.

Code:

```

1  #include "target_memory.h"
2  #include "stdio.h"
3  // Define a static lookup table for matrix memory info.
4  // These sample base addresses should match your actual
   layout.
5  static MatrixMemoryInfo matrix_memory_table[] = {
6      {'A', 0x0000},
7      {'B', 0x0100},
8      {'C', 0x0200},
9      {'V', 0x0100} // For matrix-vector multiplication.
10 };
11
12 int get_matrix_memory_table_size(void) {
13     return sizeof(matrix_memory_table) / sizeof(
        matrix_memory_table[0]);
14 }
15
16 const MatrixMemoryInfo * get_matrix_memory_info_by_index(int
   index) {
17     if (index < 0 || index >= get_matrix_memory_table_size())
18         return NULL;
19     return &matrix_memory_table[index];
20 }
21
22 void compute_memory_layout(MatrixIR *ir, MemoryLayout *layout
   ) {

```

```

23 // Compute dynamic addresses based on IR dimensions.
24 layout->base_A = 0x0000;
25 layout->base_B = layout->base_A + (ir->rows_A * ir->
    cols_A);
26 layout->base_V = layout->base_B;
27 layout->base_C = layout->base_B + (ir->rows_B * ir->
    cols_B);
28 }

```

Listing 1.7. target_memory.c

Explanation: This module calculates the base address for matrix A as 0x0000. The base address for matrix B is computed as `base_A + (rows_A * cols_A)`, and matrix C's base address follows immediately after B. For matrix-vector operations, the same base address is used for vector V as for matrix B.

2.8 print_tables.c

Purpose: Provides a utility to print the lookup tables for ISA opcodes and matrix memory base addresses. This is useful for verifying that the encoding and memory layout are correct.

Code:

```

1 #include <stdio.h>
2 #include "isa_encoding.h" // For InstructionInfo and
    lookup functions.
3 #include "target_memory.h" // For MatrixMemoryInfo and
    lookup functions.
4
5 int main() {
6     printf("=====\n");
7     printf("    Detailed Lookup Tables    \n");
8     printf("=====\\n\\n");
9
10    // Instruction Opcode Table
11    printf("=== Instruction Opcode Table ===\\n");
12    printf("%-5s %-20s %-10s\\n", "Index", "Instruction Name",
        "Opcode");
13    printf("-----\\n");
14    int instr_count = get_instruction_table_size();
15    for (int i = 0; i < instr_count; ++i) {
16        const InstructionInfo *info =
            get_instruction_info_by_index(i);
17        if (info != NULL) {
18            printf("%-5d %-20s 0x%X\\n", i, info->name, info->
                opcode);
19        }
20    }
21    printf("\\nTotal Instructions: %d\\n\\n", instr_count);

```

```

22
23 // Matrix Base Address Table
24 printf("=== Matrix Base Address Table ===\n");
25 printf("%-5s %-15s %-10s\n", "Index", "Identifier", "Base
    Address");
26 printf("-----\n");
27 int matrix_count = get_matrix_memory_table_size();
28 for (int i = 0; i < matrix_count; ++i) {
29     const MatrixMemoryInfo *info =
        get_matrix_memory_info_by_index(i);
30     if (info != NULL) {
31         printf("%-5d %-15c 0x%X\n", i, info->identifier,
            info->base_address);
32     }
33 }
34 printf("\nTotal Matrix Memory Entries: %d\n",
    matrix_count);
35
36 printf("\n=====\n");
37 return 0;
38 }

```

Listing 1.8. print_tables.c

Explanation: This program uses the lookup functions defined in `isa_encoding.c` and `target_memory.c` to print the instruction opcode table and the matrix memory table in a human-readable format.

3 Output (Screenshots and Generated Files)

This section shows the results of running the compiler. Placeholders are used here; replace them with your actual files and images.

3.1 Console Screenshots

3.2 Compiler Output File (output.isa)

When running `matrix_compiler`, an `output.isa` file is generated. Below is a snippet (placeholder) from this file:

```

1 =====
2     Generated Instruction Sequence (24-bit ISA Format)
3 =====
4
5 Processing C[0][0]
6   READ A[0][0]: 0x008000
7   READ B[0][0]: 0x00800C
8   READ A[0][1]: 0x008001
9   READ B[1][0]: 0x008011

```

```

*aditya@aditya22bail235:~/Videos/compiler_aditya$ ./matrix_compiler example.cpp
Loaded Source Code:
// OPERATION: MM MULT
// DIMENSIONS: m=3 n=4 p=5

#include <iostream>
using namespace std;

void matmul(int A[3][4], int B[4][5], int C[3][5]) {
    for (int i = 0; i < 3; i++) {
        for (int j = 0; j < 5; j++) {
            C[i][j] = 0;
            for (int k = 0; k < 4; k++) {
                C[i][j] += A[i][k] * B[k][j];
            }
        }
    }
}

int main() {
    int A[3][4] = {
        {1, 2, 3, 4},
        {5, 6, 7, 8},
        {9, 10, 11, 12}
    };

    int B[4][5] = {
        { 1, 2, 3, 4, 5},
        { 6, 7, 8, 9, 10},
        {11, 12, 13, 14, 15},
        {16, 17, 18, 19, 20}
    };

    int C[3][5] = {0};

    matmul(A, B, C);

    cout << "Result matrix C:" << endl;
    for (int i = 0; i < 3; i++) {
        for (int j = 0; j < 5; j++) {
            cout << C[i][j] << " ";
        }
        cout << endl;
    }

    return 0;
}

Parsed Info: Operation Type: Matrix-Matrix, m=3, n=4, p=5
IR: Matrix A: 3x4, Matrix B: 4x5
Optimized Info: MAC steps=60, partial mults=60
Instruction sequence generated in output.isa

```

Fig. 2. Screenshot showing successful compilation and execution of the input C++ file.


```

10 READ A[0][2]: 0x008002
11 READ B[2][0]: 0x008016
12 READ A[0][3]: 0x008003
13 READ B[3][0]: 0x00801B
14 PROG: 0x400020
15 EXE: 0x800020
16 END: 0xC00020
17
18 Processing C[0][1]
19 READ A[0][0]: 0x008000
20 READ B[0][1]: 0x00800D
21 READ A[0][1]: 0x008001
22 READ B[1][1]: 0x008012
23 READ A[0][2]: 0x008002
24 READ B[2][1]: 0x008017
25 READ A[0][3]: 0x008003
26 READ B[3][1]: 0x00801C
27 PROG: 0x400021
28 EXE: 0x800021
29 END: 0xC00021
30
31 Processing C[0][2]
32 READ A[0][0]: 0x008000
33 READ B[0][2]: 0x00800E
34 READ A[0][1]: 0x008001
35 READ B[1][2]: 0x008013
36 READ A[0][2]: 0x008002
37 READ B[2][2]: 0x008018
38 READ A[0][3]: 0x008003
39 READ B[3][2]: 0x00801D
40 PROG: 0x400022
41 EXE: 0x800022
42 END: 0xC00022
43
44 Processing C[0][3]
45 READ A[0][0]: 0x008000
46 READ B[0][3]: 0x00800F
47 READ A[0][1]: 0x008001
48 READ B[1][3]: 0x008014
49 READ A[0][2]: 0x008002
50 READ B[2][3]: 0x008019
51 READ A[0][3]: 0x008003
52 READ B[3][3]: 0x00801E
53 PROG: 0x400023
54 EXE: 0x800023
55 END: 0xC00023
56
57 Processing C[0][4]
58 READ A[0][0]: 0x008000
59 READ B[0][4]: 0x008010

```

```

60 READ A[0][1]: 0x008001
61 READ B[1][4]: 0x008015
62 READ A[0][2]: 0x008002
63 READ B[2][4]: 0x00801A
64 READ A[0][3]: 0x008003
65 READ B[3][4]: 0x00801F
66 PROG: 0x400024
67 EXE: 0x800024
68 END: 0xC00024
69
70 Processing C[1][0]
71 READ A[1][0]: 0x008004
72 READ B[0][0]: 0x00800C
73 READ A[1][1]: 0x008005
74 READ B[1][0]: 0x008011
75 READ A[1][2]: 0x008006
76 READ B[2][0]: 0x008016
77 READ A[1][3]: 0x008007
78 READ B[3][0]: 0x00801B
79 PROG: 0x400025
80 EXE: 0x800025
81 END: 0xC00025
82
83 Processing C[1][1]
84 READ A[1][0]: 0x008004
85 READ B[0][1]: 0x00800D
86 READ A[1][1]: 0x008005
87 READ B[1][1]: 0x008012
88 READ A[1][2]: 0x008006
89 READ B[2][1]: 0x008017
90 READ A[1][3]: 0x008007
91 READ B[3][1]: 0x00801C
92 PROG: 0x400026
93 EXE: 0x800026
94 END: 0xC00026
95
96 Processing C[1][2]
97 READ A[1][0]: 0x008004
98 READ B[0][2]: 0x00800E
99 READ A[1][1]: 0x008005
100 READ B[1][2]: 0x008013
101 READ A[1][2]: 0x008006
102 READ B[2][2]: 0x008018
103 READ A[1][3]: 0x008007
104 READ B[3][2]: 0x00801D
105 PROG: 0x400027
106 EXE: 0x800027
107 END: 0xC00027
108
109 Processing C[1][3]

```

```

110 READ A[1][0]: 0x008004
111 READ B[0][3]: 0x00800F
112 READ A[1][1]: 0x008005
113 READ B[1][3]: 0x008014
114 READ A[1][2]: 0x008006
115 READ B[2][3]: 0x008019
116 READ A[1][3]: 0x008007
117 READ B[3][3]: 0x00801E
118 PROG: 0x400028
119 EXE: 0x800028
120 END: 0xC00028
121
122 Processing C[1][4]
123 READ A[1][0]: 0x008004
124 READ B[0][4]: 0x008010
125 READ A[1][1]: 0x008005
126 READ B[1][4]: 0x008015
127 READ A[1][2]: 0x008006
128 READ B[2][4]: 0x00801A
129 READ A[1][3]: 0x008007
130 READ B[3][4]: 0x00801F
131 PROG: 0x400029
132 EXE: 0x800029
133 END: 0xC00029
134
135 Processing C[2][0]
136 READ A[2][0]: 0x008008
137 READ B[0][0]: 0x00800C
138 READ A[2][1]: 0x008009
139 READ B[1][0]: 0x008011
140 READ A[2][2]: 0x00800A
141 READ B[2][0]: 0x008016
142 READ A[2][3]: 0x00800B
143 READ B[3][0]: 0x00801B
144 PROG: 0x40002A
145 EXE: 0x80002A
146 END: 0xC0002A
147
148 Processing C[2][1]
149 READ A[2][0]: 0x008008
150 READ B[0][1]: 0x00800D
151 READ A[2][1]: 0x008009
152 READ B[1][1]: 0x008012
153 READ A[2][2]: 0x00800A
154 READ B[2][1]: 0x008017
155 READ A[2][3]: 0x00800B
156 READ B[3][1]: 0x00801C
157 PROG: 0x40002B
158 EXE: 0x80002B
159 END: 0xC0002B

```

```

160
161 Processing C[2][2]
162   READ A[2][0]: 0x008008
163   READ B[0][2]: 0x00800E
164   READ A[2][1]: 0x008009
165   READ B[1][2]: 0x008013
166   READ A[2][2]: 0x00800A
167   READ B[2][2]: 0x008018
168   READ A[2][3]: 0x00800B
169   READ B[3][2]: 0x00801D
170   PROG: 0x40002C
171   EXE:  0x80002C
172   END:  0xC0002C
173
174 Processing C[2][3]
175   READ A[2][0]: 0x008008
176   READ B[0][3]: 0x00800F
177   READ A[2][1]: 0x008009
178   READ B[1][3]: 0x008014
179   READ A[2][2]: 0x00800A
180   READ B[2][3]: 0x008019
181   READ A[2][3]: 0x00800B
182   READ B[3][3]: 0x00801E
183   PROG: 0x40002D
184   EXE:  0x80002D
185   END:  0xC0002D
186
187 Processing C[2][4]
188   READ A[2][0]: 0x008008
189   READ B[0][4]: 0x008010
190   READ A[2][1]: 0x008009
191   READ B[1][4]: 0x008015
192   READ A[2][2]: 0x00800A
193   READ B[2][4]: 0x00801A
194   READ A[2][3]: 0x00800B
195   READ B[3][4]: 0x00801F
196   PROG: 0x40002E
197   EXE:  0x80002E
198   END:  0xC0002E
199
200 =====
201 Total instructions generated: 165
202 Total bits: 3960
203 =====

```

Listing 1.9. Snippet of output.isa

3.3 Instruction Mapping CSV and Hex Outputs

The compiler also generates additional CSV and hexadecimal formatted outputs specifically for the PIM Compiler to facilitate detailed research analysis.

Operation	Instruction Hex	Type	Pointer	Address	Description
PROG	000000	PROG	0	0x0000	Program initialization
EXE	400000	EXE	0	0x0000	Memory operation
EXE	400004	EXE	0	0x0004	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404000	EXE	0	0x4000	Memory operation
EXE	400001	EXE	0	0x0001	Memory operation
EXE	400006	EXE	0	0x0006	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404000	EXE	0	0x4000	Memory operation
EXE	400000	EXE	0	0x0000	Memory operation
EXE	400005	EXE	0	0x0005	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404001	EXE	0	0x4001	Memory operation
EXE	400001	EXE	0	0x0001	Memory operation
EXE	400007	EXE	0	0x0007	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404001	EXE	0	0x4001	Memory operation
EXE	400002	EXE	0	0x0002	Memory operation
EXE	400004	EXE	0	0x0004	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404002	EXE	0	0x4002	Memory operation
EXE	400003	EXE	0	0x0003	Memory operation
EXE	400006	EXE	0	0x0006	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404002	EXE	0	0x4002	Memory operation
EXE	400002	EXE	0	0x0002	Memory operation
EXE	400005	EXE	0	0x0005	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404003	EXE	0	0x4003	Memory operation
EXE	400003	EXE	0	0x0003	Memory operation
EXE	400007	EXE	0	0x0007	Memory operation
EXE	410003	EXE	1	0x0003	Memory operation
EXE	404003	EXE	0	0x4003	Memory operation
END	800000	END	0	0x0000	Program termination

Hexadecimal Output

```

1 000000
2 400000
3 400004
4 410003
5 404000
6 400001
7 400006
8 410003
9 404000
10 400000
11 400005
12 410003
13 404001
14 400001
15 400007
16 410003
17 404001
18 400002
19 400004
20 410003
21 404002
22 400003
23 400006
24 410003
25 404002
26 400002
27 400005
28 410003
29 404003
30 400003
31 400007
32 410003
33 404003
34 800000

```

Listing 1.10. Hexadecimal ISA Output for PIM Compiler

3.4 C++ Input Example

Below is an excerpt from the sample C++ input file (`example.cpp`) used to drive the compilation:

```

1 // OPERATION: MM_MULT
2 // DIMENSIONS: m=3 n=4 p=5
3
4 #include <iostream>
5 using namespace std;
6
7 void matmul(int A[3][4], int B[4][5], int C[3][5]) {

```

```

8   for (int i = 0; i < 3; i++) {
9       for (int j = 0; j < 5; j++) {
10          C[i][j] = 0;
11          for (int k = 0; k < 4; k++)
12              C[i][j] += A[i][k] * B[k][j];
13      }
14  }
15  }
16
17  int main() {
18      // Matrix declarations and initialization
19      // ...
20      return 0;
21  }

```

Listing 1.11. Excerpt from example.cpp

4 Discussion and Results

The custom compiler has been tested with input files containing both matrix–matrix and matrix–vector operations. For example, using an input file with dimensions:

```
– // DIMENSIONS: m=3 n=4 p=5
```

the compiler computes the following memory layout:

- Base address for Matrix A: 0x0000
- Base address for Matrix B: $0x0000 + (3 \times 4) = 0x000C$
- Base address for Matrix C: $0x000C + (4 \times 5) = 0x0020$

The generated instruction sequence correctly encodes these addresses into 24-bit instructions, as shown in the output excerpt.

The results demonstrate:

- **Correctness:** The pipeline accurately parses, validates, and encodes the matrix multiplication operation.
- **Modularity:** Each compiler stage is separated into distinct modules, making future enhancements straightforward.
- **Flexibility:** The compiler supports both matrix–matrix and matrix–vector multiplication by adapting the memory layout and instruction generation accordingly.

5 Conclusion

The custom compiler effectively translates high-level matrix multiplication operations into a tailored 24-bit ISA instruction sequence. Detailed analysis and code segmentation confirm that each stage—from parsing through ISA encoding—is correctly implemented. Future work may involve incorporating advanced optimizations and expanding the instruction set.